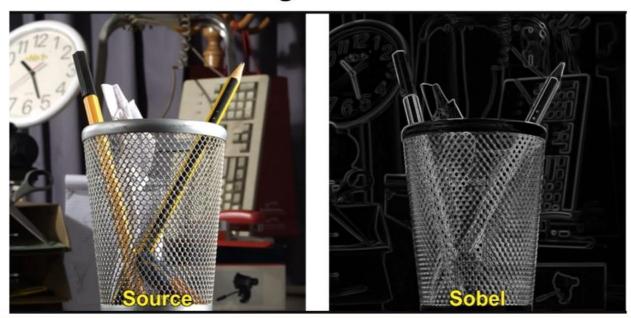


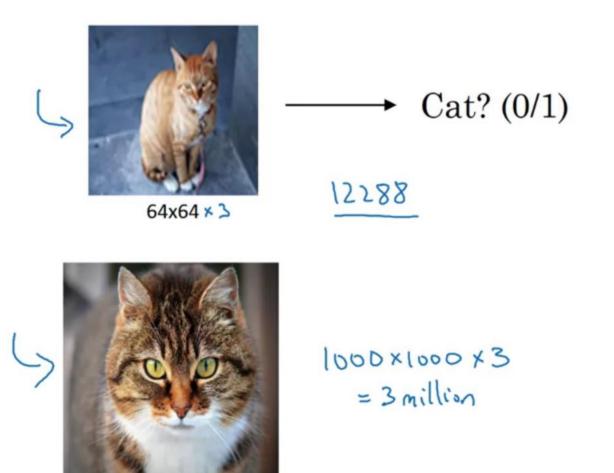
What is a CNN?

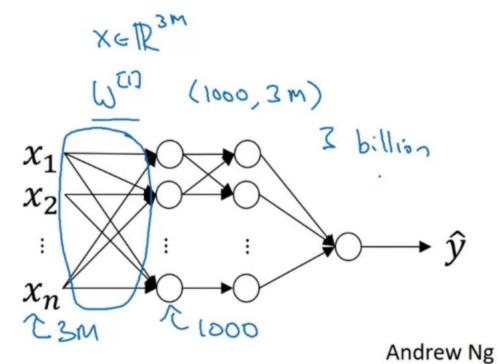
- Convolutional neural network a neural network containing convolutional layers
- Normally, convolution in Deep Learning is used for images
- An example of convolution is edge detection



Deep Learning on large images

Why CNN not regular NN?

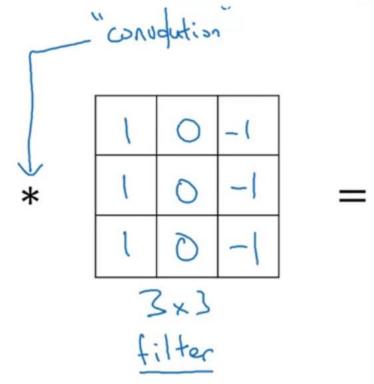


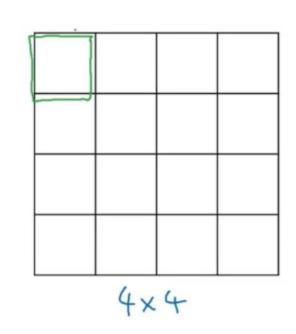


Convolution Architect: Components 1

3x1+1x1 +2x1+0x0+8x0+7x0+1x-1+8x-1+2x-1=-5

3	0 °	1	2	7	4
1	5°	8	9	3	1
2	7 °	2 ⁻¹	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9





Convolution: Elementwise Multiplication

	1					Y U
3	0	1°	2-1	7	4	"consolution"
1	5	8°		3	1	-5
2	71	2°	5	1	3	
0	1	3	1	7	8	* 1 0 -1 =
4	2	1	6	2	8	3×3
2	4	5	2	3	9	filter 4x4
		6×1	6			

Strided convolution

					_	
2 3	3 4	7 4	4	6	2	9
6 ¹	6 º	9 ²	8	7	4	3
3 -1	4 º	8 3	3	8	9	7
7	8	3	6	6	3	4
4	2	1	8	3	4	6
3	2	4	1	9	8	3
0	1	3	9	2	1	4
7×7						

3	4	4
1	0	2
-1	0	3
	3+3	

Stride = 2

Strided convolution

	X	9				
2	3	7 3	4 4	6 4	2	9
6	6	9 1	8 º	7 2	4	3
3	4	8 -1	3 º	8 3	9	7
7	8	3	6	6	3	4
4	2	1	8	3	4	6
3	2	4	1	9	8	3
0	1	3	9	2	1	4
7×7						

3	4	4				
1	0	2				
-1	0	3				
3+3						

91	ن	

Stride = 2

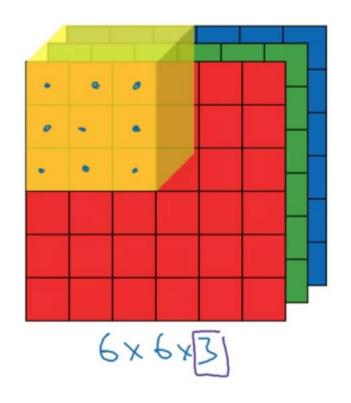
Summary of convolutions

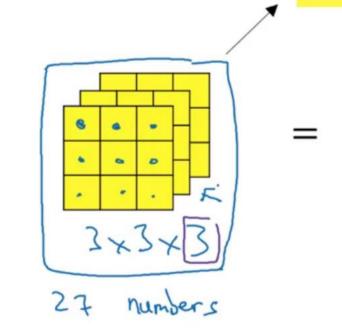
$$n \times n \text{ image}$$
 $f \times f \text{ filter}$

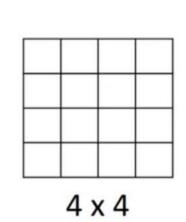
padding p stride s

$$\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor$$

Convolutions on RGB image





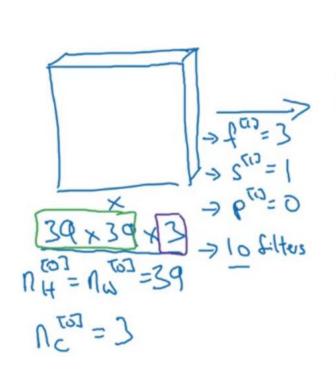


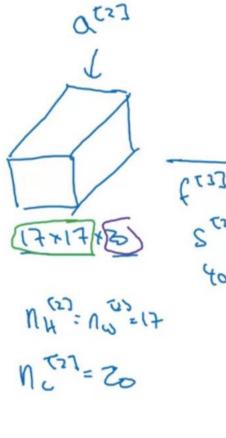
Summary of notation

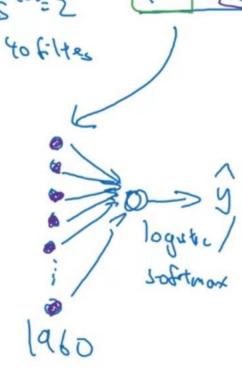
If layer 1 is a convolution layer:

```
Input: n_H \times n_W \times n_c
    f^{[l]} = filter size
                                                 Output: n tes x nw x nc cas
    p^{[l]} = padding
    s^{[l]} = \text{stride}
   n_c^{[l]} = number of filters
→ Each filter is: \( \frac{100}{2} \times \frac{100}{2} \times \left| \( \frac{100}{2} \times \left| \)
    Activations: 0 -> 1 4 × 14 × 16
                                                    ATES > M × NH + NW × NC
    Weights: fth xfth nc x nc Tli
    bias: nc - (1,11,0,41) ~ #f:(tos is layer l.
```

Example ConvNet







Types of layer in a convolutional network:

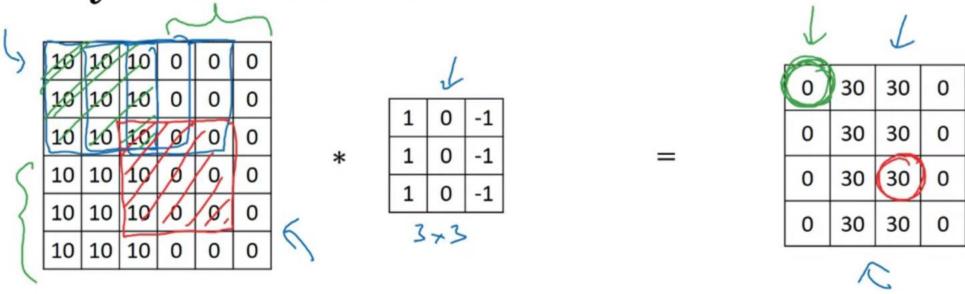
- Convolution (CONV) <
- Pooling (POOL) <
- Fully connected $(FC) \leftarrow$

$$\frac{+2p-f}{s}$$
 +1=37

Neural network example

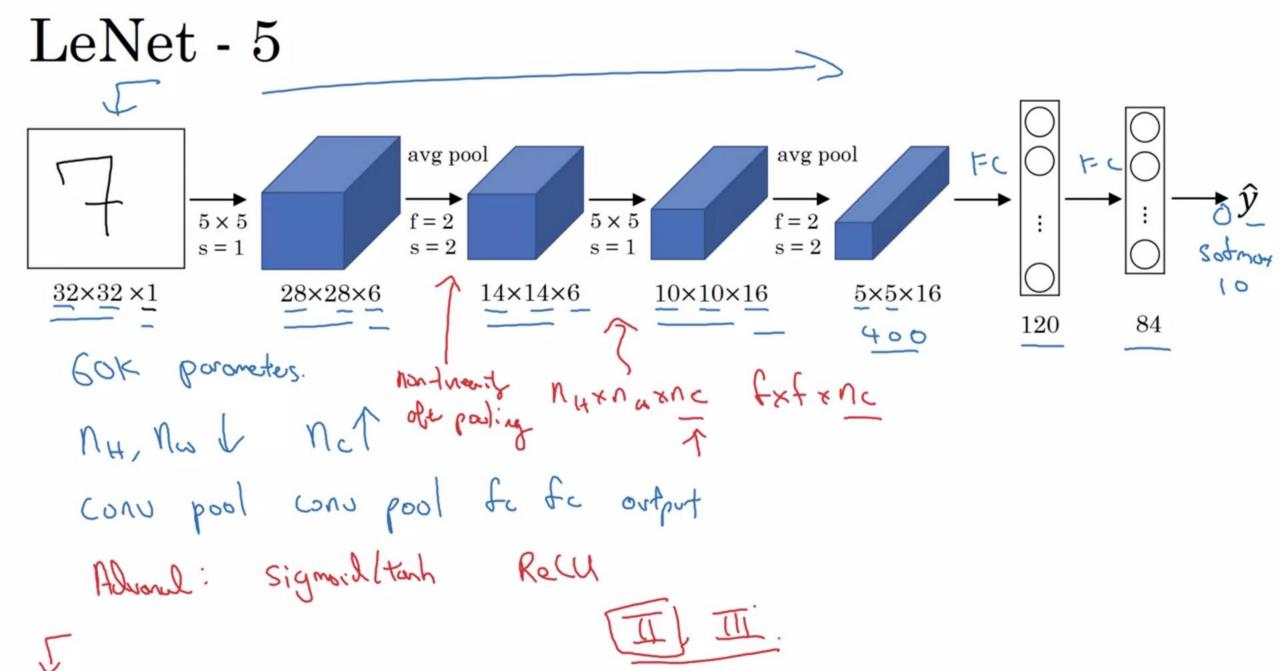
	Activation shape Activation S		# parameters
Input:	(32,32,3)	— 3,072 a ^{to]}	0
CONV1 (f=5, s=1)	(28,28,8)	6,272	208
POOL1	(14,14,8)	1,568	0
CONV2 (f=5, s=1)	(10,10,16)	1,600	416
POOL2	(5,5,16)	400	0
FC3	(120,1)	120	48,001
FC4	(84,1)	84	10,081
Softmax	(10,1)	10	841

Why convolutions

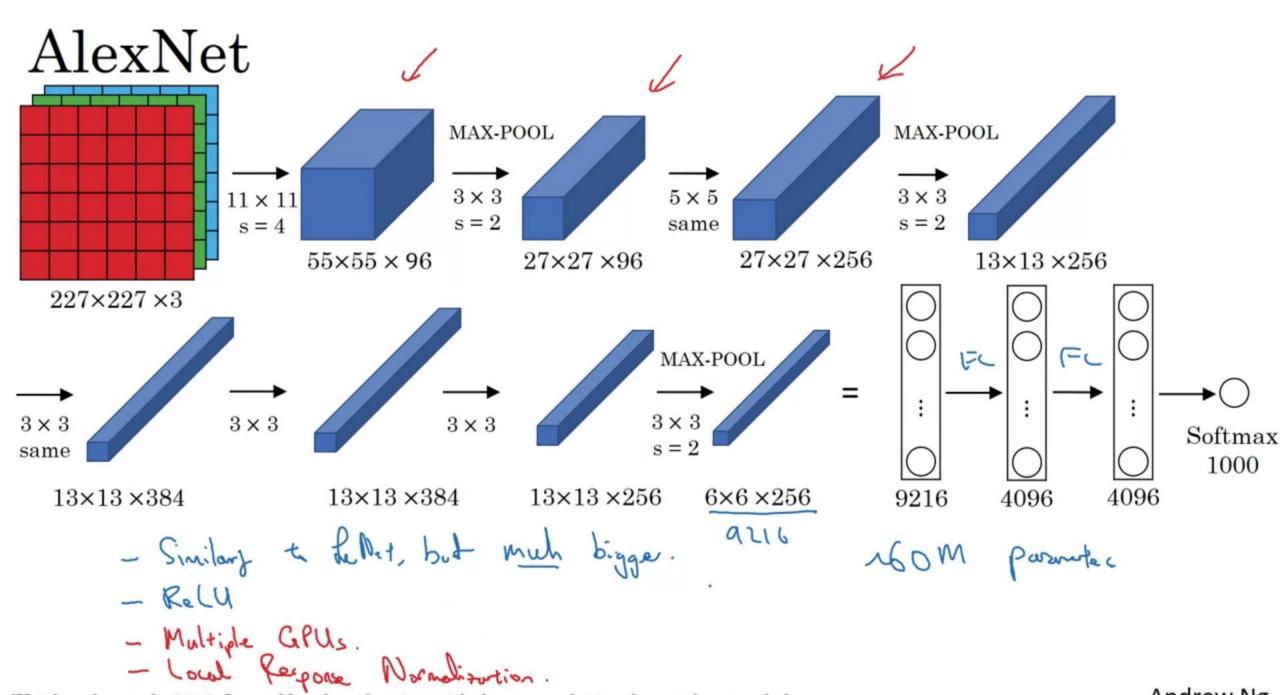


Parameter sharing: A feature detector (such as a vertical edge detector) that's useful in one part of the image is probably useful in another part of the image.

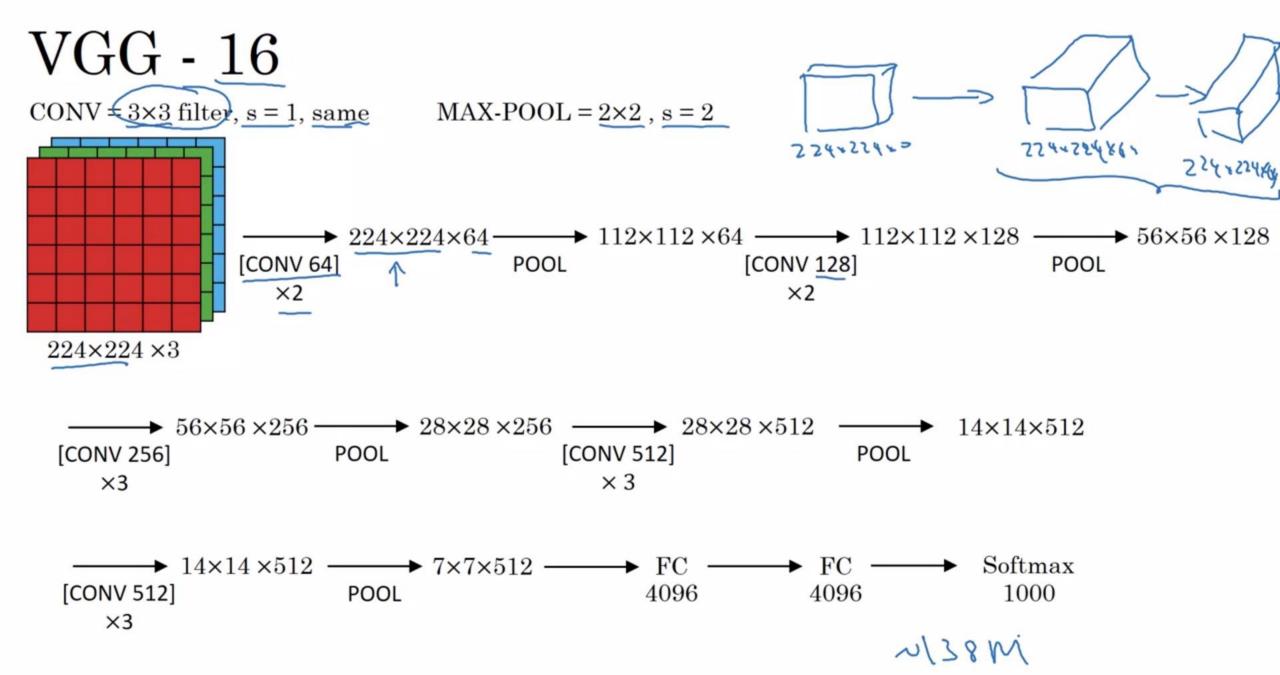
→ Sparsity of connections: In each layer, each output value depends only on a small number of inputs.



Andrew Ng

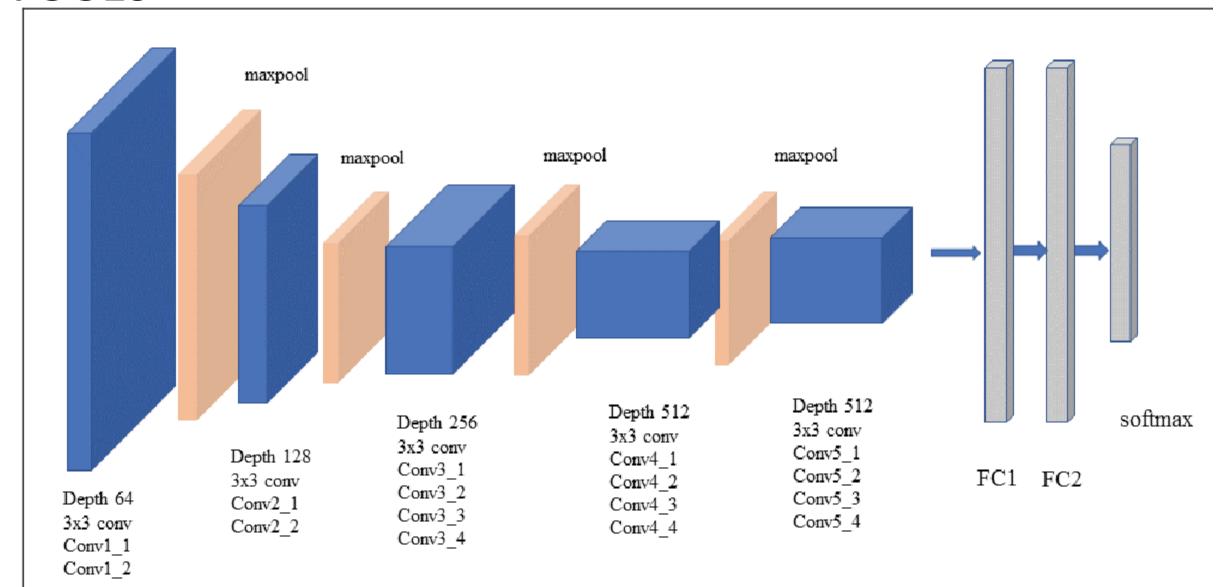


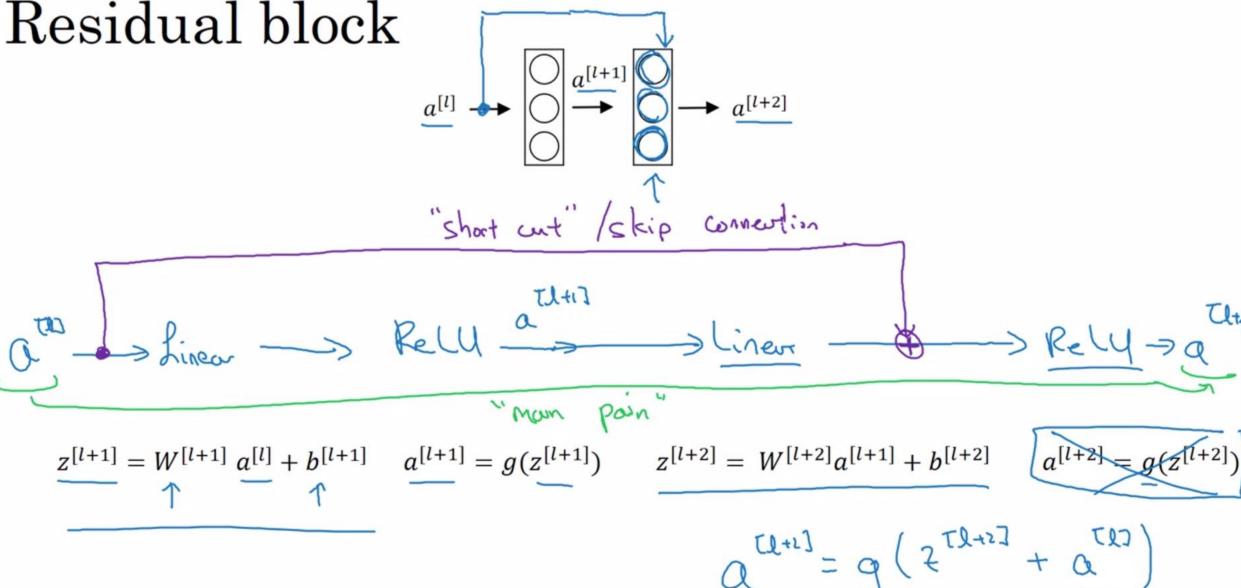
[Krizhevsky et al., 2012. ImageNet classification with deep convolutional neural networks]



Andrew Ng

VGG19





Andrew Ng

Residual Network ResNet Plain training error training error reality" theory

Andrew Ng

layers

layers

Appendix

```
# Updated to do image augmentation
train_datagen = ImageDataGenerator(
      rescale=1./255,
      rotation_range=40,
      width_shift_range=0.2,
      height_shift_range=0.2,
      shear_range=0.2,
      zoom_range=0.2,
      horizontal_flip=True,
      fill_mode='nearest')
```









Zoom



